



LLM GATEWAYS









✓ User → LLM X server Cache
500 ← Operator X X

users → request, gals → bad
user → request, gals → hack




Gateways



C Indl

Q	opus 3.5
<u> </u>	Marken 4

free → 3.5
 ↳ credits xx //



Outage


2023

Server

Rate limiting, Observability

model routing ✓

fallback



LLM costs

max

Cache



Base

user → question

9898 -



guardrails → LLM security
prompt injection ✓✓


Gateway:- ① Rate limiting

② Observability

③ model routing,
fallback

④ quadrants

⑤ cache



gateway frameworks

① Portkey → Production

↳ feature complete free

↳ paid ✓

✓

english

② Lite LLM

↳ opensource

1600+

✓

17n

③ Bigfoot → ultra fast

↳ model integration

↳ long &
low

Virtual Keys

xx

10
↑↑



gro → ---
opr → ---
chad ---

key - /

Original Key


↳ virtual key

slug



port key → fallback, ✓

✓ Cache Retry ✓




Retry → config. ✓


↳ timeout → (10s) ✓

fallback ✓

↳ change model




Fault Tolerant System



Load Balancing ?

Gen Ai

Traffic merge





chatbot




user

Ig → $\frac{10000000}{\text{sever}}$





Cache $\rightarrow$ Simple $\checkmark$
 $\rightarrow$ Semantic parse



Simple $\rightarrow$ query exactly same

Semantic $\rightarrow$ what is A_i $\leftarrow$ ✓✓

Cross encoder $\uparrow$ Tell me about A_i $\leftarrow$
attention $\rightarrow$ configs

Y not

Chatgpt

integration

→ groß url ← x

Openai

PORTKEY ✓

→ URL → Hit (↓)

→ Openrouter ✓
→ LiteLLM

Streaming !

question →

— =
(
,
,
,
,
,
)

Gateways → Enjoy

